

practical-5

April 17, 2024

1 Data Analytics II

1. Implement logistic regression using Python/R to perform classification on Social_Network_Ads.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

```
[ ]: ## Import required libraries
## import numpy as np
## import pandas as pd
## from sklearn.model_selection import train_test_split
## from sklearn.preprocessing import StandardScaler
## from sklearn.linear_model import LogisticRegression
## from sklearn.metrics import confusion_matrix, classification_report
## import matplotlib.pyplot as plt
## %matplotlib inline

## Load the dataset
## df = pd.read_csv("Social_Network_Ads.csv")
## df.head(10)

## df.info()

## df.describe()

## X = df.iloc[:, [2, 3]].values
## y = df.iloc[:, 4].values

## Split the dataset into train and test
## X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)

## Preprocessing
## Standard Scalar
## sc = StandardScaler()
```

```

# X_train = sc.fit_transform(X_train)
# X_test = sc.transform(X_test)

# classifier = LogisticRegression(random_state=0)
# classifier.fit(X_train, y_train)

# y_pred = classifier.predict(X_test)
# y_pred

# # Confusion Matrix
# cm = confusion_matrix(y_test, y_pred)
# cm

# c1_report = classification_report(y_test, y_pred)
# c1_report

# tp, fn, fp, tn = confusion_matrix(y_test, y_pred, labels=[0,1]).reshape(-1)

# print(f"Outcome Values:\nTrue Positive = {tp}\nFalse Negative = {fn}\nFalse_
↪Positive = {fp}\nTrue Negative = {tn}")

# accuracy_cm = (tp+tn)/(tp+fp+tn+fn)
# precision_cm = tp/(tp+fp)
# recall_cm = tp/(tp+fn)
# f1_score = 2/((1/recall_cm)+(1/precision_cm))

# print(f"Accuracy : {accuracy_cm}")
# print(f"Precision : {precision_cm}")
# print(f"Recall : {recall_cm}")
# print(f"F1-Score : {f1_score}")

```

```

[1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
import matplotlib.pyplot as plt
%matplotlib inline

```

```

[2]: df = pd.read_csv("Social_Network_Ads.csv")
df.head(10)

```

```

[2]:
   User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510   Male   19             19000           0
1  15810944   Male   35             20000           0
2  15668575  Female   26             43000           0

```

3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               400 non-null   int64
1   Gender                400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       400 non-null   int64
4   Purchased             400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
[4]: df.describe()
```

```
[4]:
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
[5]: X = df.iloc[:, [2, 3]].values
      y = df.iloc[:, 4].values
```

```
[6]: # Split the dataset into train and test
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
      random_state = 0)
```

```
[7]: # Preprocessing

      # Standard Scaler
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
```

```
X_test = sc.transform(X_test)
```

```
[8]: classifier = LogisticRegression(random_state=0)
classifier.fit(X_train, y_train)
```

```
[8]: LogisticRegression(random_state=0)
```

```
[9]: y_pred = classifier.predict(X_test)
y_pred
```

```
[11]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
cm
```

```
[13]: cl_report = classification_report(y_test, y_pred)
cl_report
```

```
[15]: tp, fn, fp, tn = confusion_matrix(y_test, y_pred, labels=[0,1]).reshape(-1)
```

```
[16]: print(f"Outcome Values:\nTrue Positive = {tp}\nFalse Negative = {fn}\nFalse_
      ↪Positive = {fp}\nTrue Negative = {tn}")
```

Outcome Values:

True Positive = 65

False Negative = 3

False Positive = 8

True Negative = 24

```
[17]: accuracy_cm = (tp+tn)/(tp+fp+tn+fn)
precision_cm = tp/(tp+fp)
recall_cm = tp/(tp+fn)
f1_score = 2/((1/recall_cm)+(1/precision_cm))
```

```
[18]: print(f"Accuracy : {accuracy_cm}")
print(f"Precision : {precision_cm}")
print(f"Recall : {recall_cm}")
print(f"F1-Score : {f1_score}")
```

Accuracy : 0.89

Precision : 0.8904109589041096

Recall : 0.9558823529411765

F1-Score : 0.9219858156028368

```
[ ]:
```